

OPENJUSTICE

Performance Evaluation Framework (Litigation)

Overview

This document defines the evaluation framework for assessing the performance of a legal AI tool. Each metric targets a distinct dimension of output quality, from the accuracy of legal citations to the strategic value of the analysis. Metrics are designed to be scored independently and aggregated into a composite quality score.

Evaluators should apply each metric using the guiding questions and scoring rubric provided, and writing down the answers to the evaluation questions (providing examples where possible). Scores use a consistent 1–5 scale unless otherwise noted.

Evaluators should also attempt to use a range of models (i.e. Opus vs. GPT) and detail the difference in the quality of outputs based on the questions below.

1. Citation Quality and Relevance

Definition: Measures whether the legal citations are accurate, authoritative, and factually aligned with the case at hand, and whether they conform to the required citation format.

Evaluation Questions

- **Accuracy:** Do the cited cases, statutes, and secondary sources actually exist and state what the tool claims they state?
- **Relevance:** Are the cited authorities directly applicable to the legal issues in the case? Do the facts of the cited decisions closely align with the facts of the user's scenario?
- **Authority:** Are the citations drawn from leading precedents, binding authorities, or widely recognized secondary sources? Does the tool prioritize higher-court and more recent decisions where appropriate?
- **Jurisdiction:** Are the cited authorities from the correct jurisdiction? Are out-of-jurisdiction citations flagged as persuasive rather than binding?

Score	Criteria
5 – Excellent	All citations are verified, directly on point, and drawn from leading authorities in the correct jurisdiction.
4 – Strong	Citations are accurate and mostly on point. One citation of limited relevance.
3 – Adequate	Citations are generally correct but include some tangential authorities or miss a key leading case.
2 – Weak	Several citations are poorly aligned with the facts, or from the wrong jurisdiction.
1 – Poor	Citations are largely irrelevant or unverifiable. Key authorities are missing.

2. Factual Accuracy and Hallucination Control

Definition: Measures whether the tool fabricates legal authorities, misrepresents rules or holdings, or presents unsupported assertions as established law.

Evaluation Questions

- **Existence:** Do all cited cases, statutes, and regulations actually exist? Can every reference be independently verified?
- **Representational accuracy:** Are the holdings, ratios, and rules from cited authorities accurately described? Are quotations, if used, faithful to the source?
- **Attribution:** Are all substantive legal claims supported by a cited authority? Are there assertions of law or critical arguments that lack any supporting reference?
- **Factual fidelity:** Does the tool accurately reflect the user-provided facts without distortion, omission of material details, or addition of facts not in the record?

Score	Criteria
5 – Excellent	Zero fabricated authorities. All legal claims are accurately attributed. User facts are faithfully represented.
4 – Strong	No fabricated authorities. One minor mischaracterization of a holding or minor factual imprecision.
3 – Adequate	No fabricated authorities, but some legal assertions lack supporting citations or a holding is partially misstated.
2 – Weak	One fabricated or unverifiable authority, or multiple unsupported legal assertions.
1 – Poor	Multiple fabricated authorities, significant misstatements of law, or material distortion of user-provided facts.

3. Quality of Legal Reasoning

Definition: Measures whether the tool’s analysis is logically sound, proportionate, and grounded in the cited jurisprudence without overreaching or distorting facts to fit a desired conclusion.

Evaluation Questions

- **Logical coherence:** Does the reasoning follow a clear, logical structure? Are conclusions adequately supported by the preceding analysis?
- **Proportionality:** Does the tool avoid stretching its logic to fit a preferred outcome? When instructed to favor a party (e.g., the employee), does it do so within reasonable interpretive bounds rather than distorting facts or law?
- **Jurisprudential grounding:** Is the legal reasoning anchored in the cited authorities? Can you trace each analytical step back to a supporting case, statute, or recognized legal principle?
- **Counterargument awareness:** Does the tool acknowledge weaknesses in its position or likely opposing arguments where relevant?
- **Analytical depth:** Does the analysis go beyond surface-level summary to engage meaningfully with the legal issues, or does it merely restate general legal principles?

Score	Criteria
5 – Excellent	Reasoning is rigorous, well-structured, and fully grounded in cited authorities. Proportionate advocacy with honest acknowledgment of weaknesses.
4 – Strong	Sound reasoning with clear jurisprudential support. Minor gaps in counterargument analysis or slight over-extension on one point.
3 – Adequate	Reasoning is generally logical but includes some unsupported inferential leaps or misses important nuances in the cited law.
2 – Weak	Reasoning is superficial, relies heavily on general principles without case-specific application, or noticeably stretches facts/law.
1 – Poor	Reasoning is incoherent, materially distorts facts or law, or advocates a position that cannot reasonably be supported.

4. Strategic Relevance of Analysis

Definition: Measures whether the tool identifies and prioritizes the issues that are most material to the outcome, delivering targeted, case-specific insight rather than generic legal summaries.

Evaluation Questions

- **Issue prioritization:** Does the tool identify the issues that are most critical to the case outcome? Are they prioritized the way an experienced practitioner would rank them?
- **Case specificity:** Is the analysis tailored to the user's specific factual scenario, or does it read as a generic overview that could apply to any case in the area?
- **Actionability:** Does the tool surface insights that a lawyer could directly act on (e.g., risk areas, strategic considerations, evidence gaps, procedural steps)?
- **Materiality:** Does the tool avoid raising peripheral issues that distract from the core analysis? Are all flagged concerns genuinely mission-critical?

Score	Criteria
5 – Excellent	Identifies all critical issues with correct prioritization. Analysis is case-specific, actionable, and surfaces insights a lawyer would find valuable.
4 – Strong	Covers critical issues well with mostly case-specific analysis. One minor omission or slightly generic point.
3 – Adequate	Identifies the main issues but lacks strategic depth or includes some generic content alongside case-specific analysis.
2 – Weak	Misses a critical issue or is largely generic. Limited actionable insight.
1 – Poor	Fails to identify the key legal issues or provides only boilerplate analysis unconnected to the specific facts.

5. Conciseness and Non-Redundancy

Definition: Measures whether the output is focused and efficient—free from unnecessary repetition, irrelevant content, or tangential discussion that dilutes the core analysis.

Evaluation Questions

- **Redundancy:** Does the output avoid restating the same point, authority, or argument in multiple places?
- **Relevance filtering:** Is all included content directly relevant to the legal issues at hand? Are there sections or passages that add no analytical value?
- **Brevity:** Does the tool convey its analysis in the minimum necessary space without sacrificing completeness? Does it avoid leading the user down unnecessary tangents?
- **Signal-to-noise ratio:** What proportion of the output provides substantive legal value versus filler, hedging, or generic preamble?

Score	Criteria
5 – Excellent	Every sentence adds value. No redundancy, no tangents, no filler. Tight, efficient analysis.
4 – Strong	Minimal redundancy. One or two passages could be trimmed without losing substance.
3 – Adequate	Some noticeable repetition or tangential content, but the core analysis remains clear.
2 – Weak	Significant redundancy or substantial irrelevant content that obscures the key analysis.
1 – Poor	Highly repetitive or padded with generic content. Core analysis is difficult to extract.

6. Document Structure and Formatting

Definition: Measures whether the output follows professional legal document conventions, with clear organization, proper heading hierarchy, and correct citation formatting that supports readability.

Evaluation Questions

- **Heading structure:** Does the document use clear, descriptive headings? Is the heading hierarchy logical and consistent (e.g., headers nest properly under parent sections)?
- **Citation formatting:** Are in-text citations formatted per the McGill Guide (or applicable standard) throughout, including pinpoints, short forms, and secondary sources?
- **Professional conventions:** Does the format align with what would be acceptable in common legal documents (e.g., memoranda, opinions, briefs)?
- **Readability:** Does the formatting enhance comprehension? Are sections logically sequenced? Is the document easy to scan and navigate?
- **Format:** Do all citations conform to the McGill Guide (or the applicable citation standard)? Are pinpoint references included where appropriate?

Score	Criteria
5 – Excellent	Professional-grade formatting. Consistent heading hierarchy, correct citation format throughout, and a structure that mirrors standard legal documents.

4 – Strong	Well-structured with minor formatting inconsistencies (e.g., one citation format error or slight heading irregularity).
3 – Adequate	Acceptable structure but noticeable formatting issues: inconsistent headings, several citation errors, or unclear section flow.
2 – Weak	Poor organization that hampers comprehension. Significant formatting or citation errors throughout.
1 – Poor	No discernible structure. Formatting does not meet professional legal standards.

7. Completeness of Legal Coverage

Definition: Measures whether the tool addresses all relevant legal issues, applicable statutes, and potential causes of action raised by the facts without omitting material areas of law.

Evaluation Questions

- **Issue spotting:** Does the tool identify all significant legal issues that a competent lawyer would flag given the facts?
- **Statutory and regulatory coverage:** Are all relevant statutes, regulations, and legislative provisions addressed?
- **Multi-issue handling:** In complex scenarios involving multiple causes of action or overlapping areas of law, does the tool address each adequately?

Score	Criteria
5 – Excellent	All material legal issues, statutes, and causes of action are identified and addressed with appropriate depth.
4 – Strong	Covers nearly all issues. One minor area of law omitted or given insufficient treatment.
3 – Adequate	Covers the primary issues but misses a secondary legal issue or applicable jurisprudence.
2 – Weak	Misses one or more significant legal issues or entire areas of applicable law.
1 – Poor	Grossly incomplete. Major legal issues or causes of action are absent.

8. Consistency and Internal Coherence

Definition: Measures whether the output is internally consistent free from contradictory statements, conflicting legal positions, or logical inconsistencies across sections.

Evaluation Questions

- **Cross-section consistency:** Do different sections of the output align with one another? Are there contradictory conclusions or conflicting legal positions?
- **Factual consistency:** Are the user's facts represented uniformly throughout, without drift or variation?

- **Terminological consistency:** Are legal terms, party names, and case references used consistently throughout the document?

Score	Criteria
5 – Excellent	Fully consistent across all sections. No contradictions, factual drift, or terminological inconsistencies.
4 – Strong	Essentially consistent with one minor terminological or factual inconsistency.
3 – Adequate	Generally consistent but contains a noticeable contradiction or inconsistency that may confuse the reader.
2 – Weak	Multiple inconsistencies that undermine the reliability of the analysis.
1 – Poor	Major internal contradictions that render sections of the output unreliable.

9. Appropriate Tone and Professional Register

Definition: Measures whether the output maintains the professional tone, objective register, and appropriate confidence level expected of legal analysis.

Evaluation Questions

- **Objectivity:** Does the analysis maintain an appropriate level of objectivity, avoiding inflammatory or overly casual language?
- **Confidence calibration:** Does the tool appropriately hedge uncertain positions and express stronger confidence only where the law clearly supports it?
- **Audience fit:** Is the tone appropriate for the intended audience (e.g., client-facing vs. internal memo vs. court filing)?

Score	Criteria
5 – Excellent	Consistently professional tone. Confidence levels accurately reflect the strength of the underlying legal support.
4 – Strong	Professional tone throughout. Minor calibration issues (e.g., slightly overconfident on one uncertain point).
3 – Adequate	Generally professional but with noticeable tonal issues: overly hedged, excessively confident, or inconsistently formal.
2 – Weak	Tone is frequently inappropriate for legal analysis (too casual, too aggressive, or poorly calibrated).
1 – Poor	Unprofessional register that would be unacceptable in a legal setting.